

DMBLOCK Assignment 1

Lukáš Častven

Slovenská technická univerzita v Bratislave
Fakulta informatiky a informačných technológií
xcastven@stuba.sk

17 march 2024

Contents

1	Programming language choice	2
1.1	Environment details	2
2	Phase 1 - Simple Coin	2
2.1	Assignment	2
2.2	Implementation	2
2.3	Tests	6
3	Phase 2 - Trust and Consensus	7
3.1	Assignment	7
3.2	Implementation	8
3.3	Tests	9

1 Programming language choice

For assignment we received a set of Java source codes, which are to be used to implement given tasks. Since there is a 5 point bonus for rewriting these source codes and implementing this assignment in different language, I choose to do it in Rust [1], because I wanted to better myself in it.

I rewrote provided source codes for each of the three phases into Rust. My source codes roughly match the provided ones.

Each phase is in it's separate directory with test as listed in the assignment.

1.1 Environment details

This assignment was developed and tested in this environment:

- OS: Arch Linux, Kernel 6.6.21-1-lts,
- Rust compiler: rustc 1.76.0 (07dca489a 2024-02-04),
- Cargo: cargo 1.76.0 (c84b36747 2024-01-18),
- Libraries: each version is listed in corresponging Cargo.toml file.

2 Phase 1 - Simple Coin

2.1 Assignment

Centralized authority FIIT accepts transactions from users. Implement logic for processing transactions into a ledger. FIIT groups transactions into into a pseudo blockchain. In each block, FIIT will receive list of transactions which are validated, applied and subset of valid ones is returned.

These transactions can depend on each other, there may be double-spends, and otherwise invalid transactions.

In this phase, implement a UTXO FIIT chain, which in each epoch takes a list of proposed transactions, validates them, applies valid ones to its internal state and returns a subset of valid ones. The size of the subset isn't defined.

Also implement a version in which, not only validity is checked, but also transactions are processed in order to maximize received fees.

2.2 Implementation

Implementation for this phase is in submodule *fitcoin*. This submodule is a Rust library crate, so I could easily reuse it in next phases with importing it as a dependency.

Required handlers are in rust file *fitcoin/src/handler.rs*. Both implement common trait `TxHandler`:

Listing 1: TxHandler

```

1 pub trait TxHandler<'a> {
2     /// Each epoch accepts unordered vector of proposed transactions.
3     /// Checks validity of each, internally updates the UTXO pool, and
4     /// returns vector of valid ones.
5     ///
6     /// # Beware
7     /// Transactions can be dependent on other ones. Also, multiple
8     /// transactions can reference same output.
9     fn handle(&mut self, possible_txs: Vec<&'a Tx>) -> Vec<&'a Tx>;
10
11     /// Returns reference to internal pool
12     fn pool(&self) -> &UTXOPool;
13
14     /// Returns mutable reference to internal pool
15     fn pool_mut(&mut self) -> &mut UTXOPool;
16
17     /// Moves internal pool, while consuming self
18     fn move_pool(self) -> UTXOPool;
19
20     /// Checks if:
21     ///     1. All UTXO inputs are in pool
22     ///     2. Signatures on inputs are valid
23     ///     3. No UTXO is used more than once
24     ///     4. Sum of outputs is not negative
25     ///     5. Sum of inputs >= Sum of outputs
26     fn is_tx_valid(&self, tx: &Tx) -> bool;
27
28     /// Filters independent txs from dependent ones,
29     /// applies them and returns both sets
30     fn handle_independent(
31         &mut self,
32         txs: Vec<&'a Tx>
33     ) -> (Vec<&'a Tx>, Vec<&'a Tx>);
34
35     /// Applies given tx to the internal pool
36     fn apply_tx(&mut self, tx: &Tx);
37
38
39     fn is_input_in_pool(&self, input: &Input) -> bool;
40 }

```

Methods `is_tx_valid`, `handle_independent`, `apply_tx`, `is_input_in_pool` have default implementations in the trait declaration, because all of them are used in both handlers. Methods `handle`,

pool, pool_mut, and move_pool are implemented in each handler.

The `is_tx_valid` function is self explanatory, and all checks, which are performed are listed in its doc comment.

The `handle_independent` function is used to separate transactions, which are not dependent on transactions in the currently proposed list and can be applied right away, from invalid and dependent transactions. The valid independent ones are also applied to the internal state.

Listing 2: `handle_independent`

```
1 /// Filters independent txs from dependent ones,
2 /// applies them and returns both sets
3 fn handle_independent(
4     &mut self,
5     txs: Vec<&'a Tx>
6 ) -> (Vec<&'a Tx>, Vec<&'a Tx>) {
7     let mut handled = vec![];
8     let mut dependent = vec![];
9     let tx_set: HashSet<u8; 32> = txs.iter().map(|tx| tx.hash()).collect();
10
11     for &tx in txs.iter() {
12         if tx.inputs().iter().all(|i| self.is_input_in_pool(i)) {
13             // tx is only dependent on outputs in pool
14             if self.is_tx_valid(tx) {
15                 self.apply_tx(tx);
16                 handled.push(tx);
17             }
18         } else if tx
19             .inputs()
20             .iter()
21             .any(|i| tx_set.contains(&i.output_tx_hash()))
22         {
23             // tx is dependent on some outputs from this batch
24             dependent.push(tx)
25         }
26     }
27
28     (handled, dependent)
29 }
```

The `apply_tx` function takes a valid transaction and mutates internal state according to the transaction.

Listing 3: `apply_tx`

```
1 /// Applies given tx to the internal pool
2 fn apply_tx(&mut self, tx: &Tx) {
3     for input in tx.inputs().iter() {
4         self.pool_mut().remove_utxo(&input_to_utxo(input));
5     }
6     for (i, output) in tx.outputs().iter().enumerate() {
7         let utxo = UTXO::new(tx.hash(), i.try_into().unwrap());
8         self.pool_mut().add_utxo(utxo, &output)
9     }
10 }
```

Handler

Handler implements the FIITcoin chain logic of validating, applying and returning valid subset of transactions. It's functionality is very simple, loop until there are no dependent transactions, which are retrieved with `handle_independent` function. At the end returns the subset of valid, applied transactions.

Listing 4: `Handler::handle`

```
1 fn handle(&mut self, possible_txs: Vec<&'a Tx>) -> Vec<&'a Tx> {
2     let mut handled: Vec<&'a Tx> = vec![];
3     let mut to_handle = possible_txs;
4
5     loop {
6         let (independent, dependent) = self.handle_independent(to_handle);
7         handled.extend(independent);
8         if dependent.is_empty() {
9             break;
10        }
11        to_handle = dependent;
12    }
13
14    handled
15 }
```

MaxFeeHandler

This implementation of a handler processes transactions in order to maximize collected fees. Since there is no maximum count of processed transactions (if there was the problem would be NP

hard and very similar to Knapsack Problem [2]), I choose very simple heuristic how to achieve this... Process all valid transactions.

Firstly, for each transaction a fee is calculated. The fee is a difference between sum of input values and sum of output values.

$$Fee = \sum_{n=1}^{|inputs|} inputs_n - \sum_{n=1}^{|outputs|} outputs_n$$

Then the initial list of proposed transactions is sorted by their fee, from highest to lowest. And then are processed just like in 2.2.

Listing 5: MaxFeeHandler::handle

```

1 fn handle(&mut self, possible_txs: Vec<&'a Tx>) -> Vec<&'a Tx> {
2     let tx_map: HashMap<u8; 32>, &'a Tx> =
3         possible_txs.iter().map(|&tx| (tx.hash(), tx)).collect();
4
5     let mut with_fees: Vec<(u64, &Tx)> = possible_txs
6         .iter()
7         .filter_map(|&tx| match self.calc_fee(tx, &tx_map) {
8             Some(fee) => Some((fee, tx)),
9             None => None,
10        })
11        .collect();
12    with_fees.sort_unstable_by(|tx1, tx2| tx1.0.cmp(&tx2.0));
13    with_fees.reverse();
14
15    let mut handled: Vec<&'a Tx> = vec![];
16    let mut to_handle = with_fees.iter().map(|tx| tx.1).collect();
17
18    loop {
19        let (independent, dependent) = self.handle_independent(to_handle);
20        handled.extend(independent);
21        if dependent.is_empty() {
22            break;
23        }
24        to_handle = dependent;
25    }
26
27    handled
28 }
```

2.3 Tests

Tests for this phase are in *fiitcoin/tests*, and are implemented according to provided test names in assignment.

Test 1 through 7 are in *fitcoin/tests/is_tx_valid_test.rs*. However test case number 7 isn't implemented. This is because in the provided source codes, Java's `double` is used, which is a signed representation of a fractional number [3]. I used Rust's `u32` [4] to denominate value of an output, thus testing if a sum of all outputs is negative is meaningless. Even if I created a raw bytes representation of a transaction, and instead of unsigned integer encoded a signed one, it would be treated as a large unsigned integer, in which case the check for sum of inputs being less then or equal to sum of outputs would fail, thus creating the same result.

Test 8 through 15 are in *fitcoin/tests/handler_test.rs*. And max fee tests 1 through 3 are in *fitcoin/tests/max_fee_handler_test.rs*.

To run these tests, navigate to *fitcoin* and run this command:

```
$ RUST_LOG=debug cargo test --release
```

3 Phase 2 - Trust and Consensus

3.1 Assignment

Implement algorithm for distributed consensus based on relationship graph. Network is an oriented graph, in which each edge represents a trust relationship between two nodes. If there is an edge $A \rightarrow B$ it means that node B is a follower of node A (A is a followee of B) and listens to transactions proposed by A .

Nodes in network are either trusted or byzantine. Each trusted node should reach consensus with other peers in the network. Implement trusted node, which defines how each trusted node in network behaves. Then test this network in a simulation with different relationship graphs and with different parameters. At the end of the simulation, each node should return the same subset of transactions, upon which consensus was reached. Assume all transactions are valid. Different simulation parameters:

- Probability of an edge existing = $p_{graph} \in \{0.1, 0.2, 0.3\}$,
- Probability of an node being byzantine = $p_{byzantine} \in \{0.15, 0.3, 0.45\}$,
- Probability of a transaction distribution to node = $p_{tx_dist} \in \{0.01, 0.05, 0.1\}$,
- Number of rounds in simulation = $rounds \in \{10, 20\}$.

Also implement a byzantine node. This node acts like an adversary in the network, and is trying to disrupt the network. There can up to 45% of byzantine nodes in the network. Their behaviour is:

- Dead - don't resend any transactions,
- Selfish - only resend their transactions,
- Mix - switch between previous two behaviours.

At the end of the simulation, all trusted nodes must return the same subset of transactions. The size of this subset should be maximal, and time to reach consensus should be reasonable.

3.2 Implementation

Implementation for this phase is in submodule *consensus*. This submodule is also a Rust library crate, because I didn't know if I will reuse it later or not.

Implementations of trusted and byzantine node are in *consensus/src/node.rs*, and both of them share common trait *Node*:

Listing 6: Node trait

```
1 pub trait Node<const N: usize> {
2     /// If 'ith' entry is 'true' then this Node follows the 'ith' Node
3     fn followees_set(&mut self, followees: [bool; N]);
4
5     /// Initializes proposed set of txs
6     fn pending_txs_set(&mut self, pending_txs: HashSet<Tx>);
7
8     /// Returns proposed txs, which shall be send to this Node's followers.
9     /// After final round, the behaviour changes and it will return txs,
10    /// on which consensus was reached.
11    fn followers_send(&self) -> &HashSet<Tx>;
12
13    /// Candites from different Nodes
14    fn followees_receive(&mut self, candidates: &Vec<Candidate>);
15
16    fn is_byzantine(&self) -> bool;
17 }
```

TrustedNode

Trusted node starts with transmitting its transactions, and after receiving a set of proposed transactions from followees it starts to track how many distinct peers proposed each transaction, and adds each proposed transaction to the ones it will transmit in next round.

A trusted node will assume a consensus was reached upon a transaction, if certain number of peers propose it back to it. The number is called `consensus_threshold` in code and is calculated with this formula:

$$CT = \min(1, Probable\textit{followers} - Probable\textit{byzantinenodes})$$

Each node has access to simulation's parameters so it can compute how many followers it should have and how many byzantine nodes there are in network.

At the end of a simulation, each node returns the set of transactions upon which consensus was reached on.

ByzantineNode

Byzantine node is not interesting. It either doesn't resend any transactions, only resends its transactions and no other, or switches between these two behaviours.

3.3 Tests

There is only one test, and that is in *consensus/tests/simulation.rs*, in which for all permutations of the simulation parameters are parallelly tested. Each simulation has a result, and this result is written to */tmp/sim-result.txt* file (in *consensus*, there is a file from this test that I ran).

Since randomness is used to initialize each simulation, if it fails it is rerun, but not more than 3 times, in order to make sure that it failed due to a bug, not due to an edge case.

Each line in the *consensus/sim-result.txt* contains what parameters were used, how long the initialization of the simulation took, what seeds were used for randomness, how long the simulation took and what was the size of the transaction subset upon which consensus was reached on. If the simulation fails more than 3 times, there would be an additional entry in the line, which mentions this, but I didn't encounter any.

Test cases in assignment are covered by this one simulation which uses all possible permutations of parameters.

To run these tests, navigate to *consensus* and run this command:

```
RUST_LOG=info cargo test --release
```

References

- [1] "Rust Programming Language — rust-lang.org." <https://www.rust-lang.org/>. [Accessed 17-03-2024].
- [2] Wikipedia, "Knapsack problem — Wikipedia, the free encyclopedia." <http://en.wikipedia.org/w/index.php?title=Knapsack%20problem&oldid=1212701319>, 2024. [Online; accessed 17-March-2024].
- [3] "DOUBLE PRECISION data type — docs.oracle.com." <https://docs.oracle.com/javadb/10.10.1.2/ref/rrefsqljdoubleprecision.html>. [Accessed 17-03-2024].
- [4] "u32 - Rust — doc.rust-lang.org." <https://doc.rust-lang.org/std/primitive.u32.html>. [Accessed 17-03-2024].